

# Multi-Agent Insurance System

Course Name: Agentic AI

**Institution Name:** Medicaps University – Datagami Skill Based Course

*Student Name(s) & Enrolment Number(s):*

Sr no	Student Name	Enrolment Number
1.	Nandini Sharma	EN22CS301635
2.	Mustafa Shaikh	EN22CS301624
3.	Nishkarsh Sharma	EN22CS301659
4.	Naman Mehta	EN22CS301629
5.	Naman Raghuvanshi	EN22CS301632

*Group Name: 09D6*

*Project Number: AAI-20*

*Industry Mentor Name:*

*University Mentor Name: Prof. Nishant Shrivastav*

*Academic Year: January-June 2026*

## Problem Statement & Objectives

### 1. Problem Statement

The insurance industry deals with highly complex policy documents, regulatory clauses, and compliance requirements. Manual analysis of policies is:

- Time-consuming
- Error-prone
- Difficult to scale
- Susceptible to inconsistency

Traditional AI systems struggle with **context overflow**, hallucinations, and lack of traceability when processing large insurance documents.

There is a need for a **collaborative, specialized multi-agent system** that can automate policy research and structured report generation while maintaining accuracy, compliance, and data security.

### 2. Project Objectives

The main objectives of this project are:

- Design and implement a **Multi-Agent Insurance System**
- Separate responsibilities between specialized agents
- Automate insurance policy research and reporting
- Ensure structured and traceable outputs
- Implement secure local LLM processing
- Maintain workflow state management
- Reduce hallucinations through structured context passing
- Provide a scalable and modular architecture

### 3. Scope of the Project

The project scope includes:

- Development of a **two-agent collaborative system**
- Integration of:
  - LLM for policy analysis
  - External API for data retrieval
  - MongoDB for persistence
- REST API-based interaction
- Sequential task orchestration using agent framework
- Error handling, retry mechanisms, and logging
- Structured output generation for insurance policy reports

### Proposed Solution

The system implements a **Collaborative Multi-Agent Architecture** where tasks are divided between:

- Researcher Agent (Fact extraction & evidence grounding)
- Writer Agent (Synthesis & structured documentation)

Orchestration is handled using **CrewAI**, ensuring sequential execution and automatic context passing.

LLM inference is handled locally using **Ollama** to ensure privacy and PII protection.

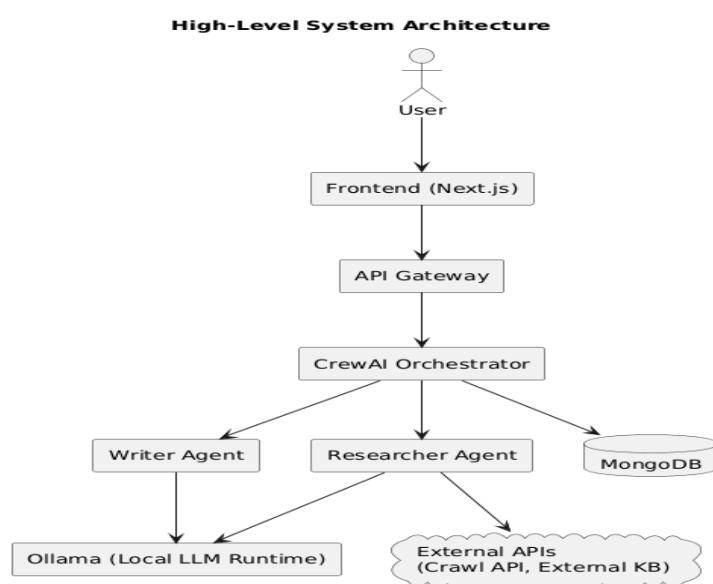
Data persistence is implemented using **MongoDB**.

## 1. Key Features

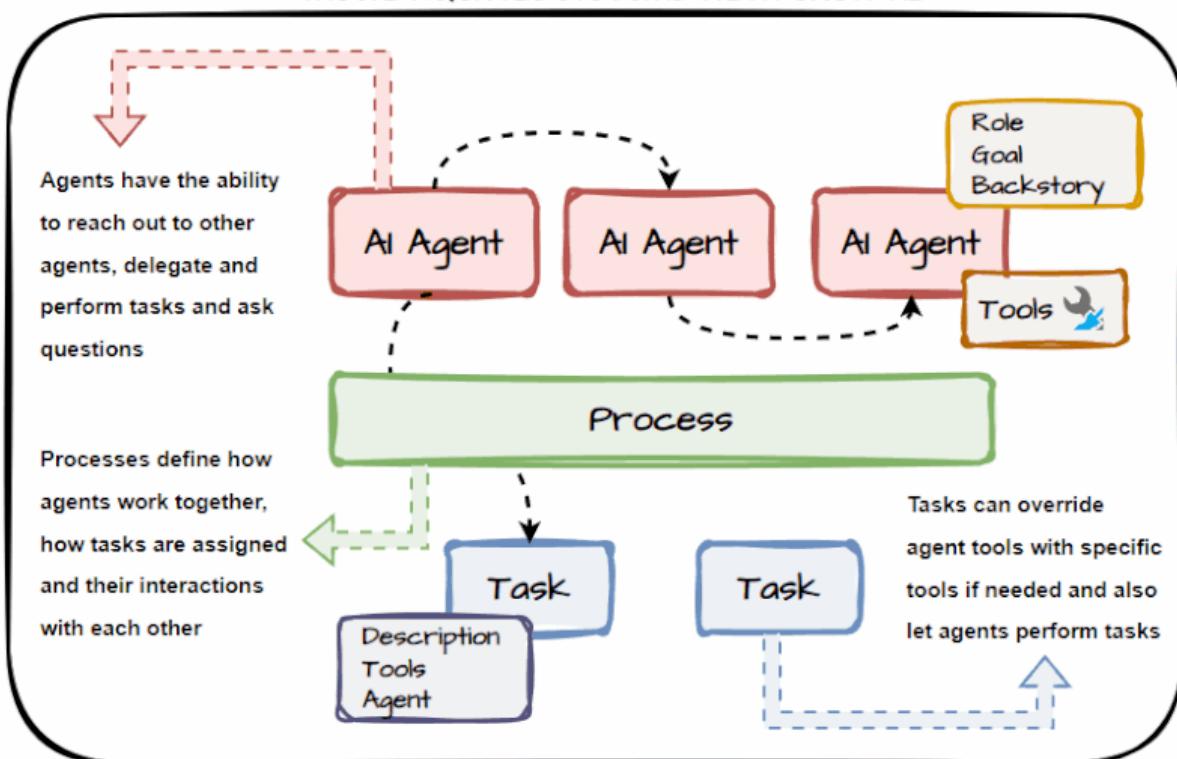
- Agent Specialization (Researcher & Writer)
- Sequential Task Orchestration
- Automatic Context Passing
- Structured JSON Evidence Generation
- RAG-based Information Retrieval
- Local LLM Processing (Privacy-first)
- Workflow State Management
- Retry with Exponential Backoff
- Circuit Breaker Pattern
- Semantic Caching
- REST API Endpoints
- Structured Report Generation

## 2. Overall Architecture / Workflow

### High-Level Architecture



### MULTI-AGENTIC SYSTEMS WITH CREW AI



### System Layers

#### Client Layer

User submits insurance query

#### API Layer

Receives request and triggers workflow

#### Orchestration Layer

- Managed by CrewAI
- Handles task sequencing
- Transfers structured context

#### Agent Layer

- Researcher Agent
- Writer Agent

#### Data Layer

- MongoDB collections:

- workflow\_states
- research\_data
- final\_outputs

### **External Services**

- Ollama (LLM inference)
- Crawl API

### **Workflow Execution Steps**

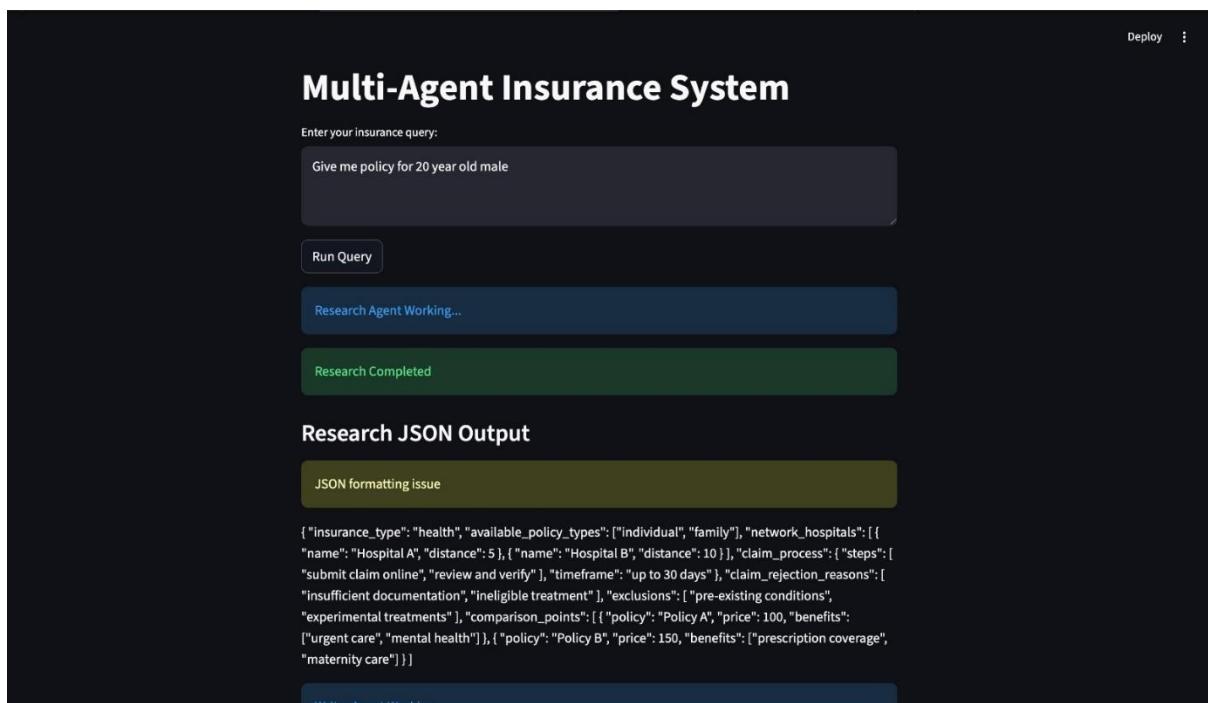
1. User submits insurance query
2. API creates workflow ID
3. Researcher Agent:
  - Queries LLM
  - Fetches external data
  - Validates findings
  - Generates structured JSON
4. Writer Agent:
  - Receives structured context
  - Synthesizes report
  - Applies formatting & compliance
5. Final output stored in database
6. Status updated to “Completed”
7. Structured response returned to user

### 3. Tools & Technologies Used

Layer	Technology
Agent Orchestration	CrewAI
LLM Integration	LangChain
LLM Runtime	Ollama
Database	MongoDB
Backend	FastAPI / Flask
Programming Language	Python 3.9+
External Data	Crawl API
Optional Caching	Redis
Logging	OpenTelemetry
Security	IAM Integration

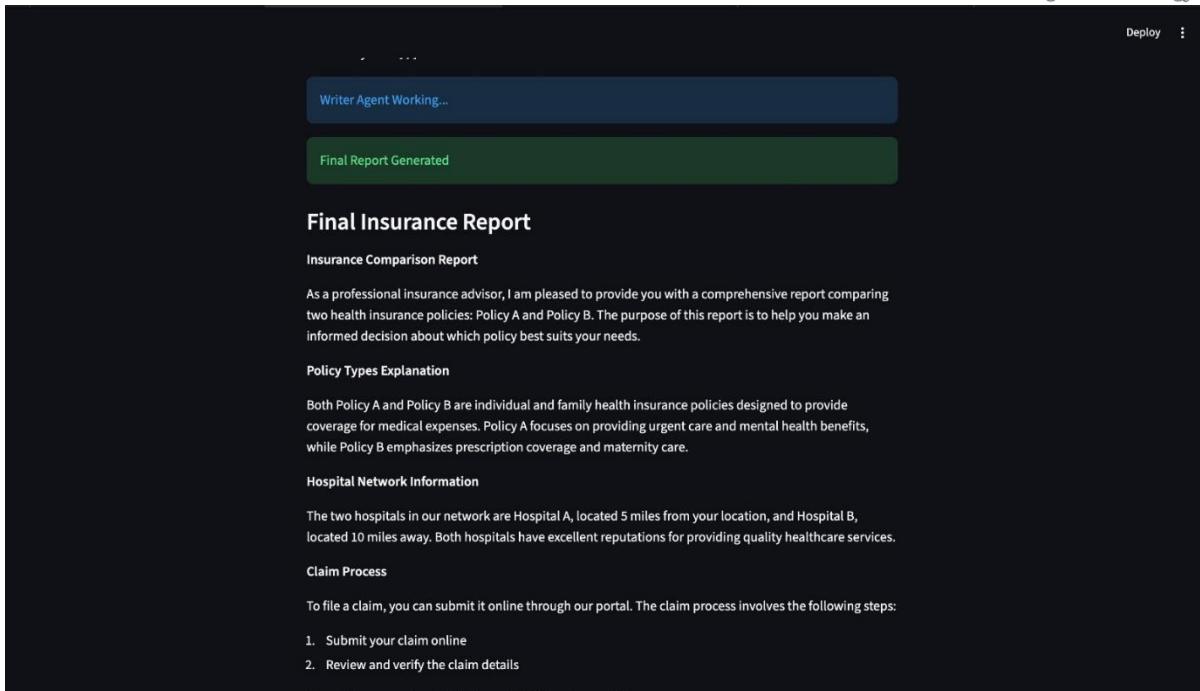
## Results & Output

### 1. Screenshots / Outputs



The screenshot shows a user interface for a "Multi-Agent Insurance System". At the top, there is a dark header bar with a "Deploy" button and a three-dot menu icon. Below the header, the main title "Multi-Agent Insurance System" is displayed in white. A text input field contains the query "Give me policy for 20 year old male". Below the input field is a "Run Query" button. The system is currently processing the query, indicated by a blue progress bar labeled "Research Agent Working...". Once the process is completed, a green bar indicates "Research Completed". Below the progress bars, the title "Research JSON Output" is shown, followed by a yellow warning box stating "JSON formatting issue". A large text area displays the JSON response:

```
{ "insurance_type": "health", "available_policy_types": ["individual", "family"], "network_hospitals": [ { "name": "Hospital A", "distance": 5 }, { "name": "Hospital B", "distance": 10 } ], "claim_process": { "steps": [ "submit claim online", "review and verify" ], "timeframe": "up to 30 days" }, "claim_rejection_reasons": [ "insufficient documentation", "ineligible treatment" ], "exclusions": [ "pre-existing conditions", "experimental treatments" ], "comparison_points": [ { "policy": "Policy A", "price": 100, "benefits": [ "urgent care", "mental health" ] }, { "policy": "Policy B", "price": 150, "benefits": [ "prescription coverage", "maternity care" ] } ] }
```



The screenshot shows a user interface for generating a final insurance report. At the top, there is a blue progress bar with the text "Writer Agent Working...". Below it, a green bar indicates "Final Report Generated". The main content area is titled "Final Insurance Report" and includes sections for "Insurance Comparison Report", "Policy Types Explanation", "Hospital Network Information", and "Claim Process". Each section contains descriptive text about its respective topic.

The system successfully generates:

- Structured JSON research outputs
- Formatted insurance coverage reports
- Workflow status logs
- Stored research evidence with traceability

## 2. Reports / Dashboards / Models

The system produces:

- Policy Coverage Summary Reports
- Structured JSON Output Files
- Research Confidence Scores
- Workflow Logs
- Quality Metrics:
  - Completeness Score
  - Accuracy Score
  - Readability Score

MongoDB stores:

- Research evidence
- Final formatted outputs
- Workflow state transitions

### **3. Key Outcomes**

- Successfully implemented a **collaborative multi-agent system**
- Reduced hallucination through task isolation
- Achieved structured, traceable outputs
- Ensured PII-safe processing using local LLM
- Built scalable modular architecture
- Implemented error handling & retry strategy
- Enabled state persistence and resume capability
- Demonstrated practical application of Agentic AI in Insurance

## Conclusion

The Multi-Agent Insurance System demonstrates how **Agentic AI architectures** can transform complex insurance workflows.

By dividing responsibilities between:

- Researcher Agent (Fact extraction)
- Writer Agent (Synthesis & formatting)

the system significantly improves:

- Accuracy
- Traceability
- Modularity
- Scalability
- Compliance

The use of CrewAI for orchestration, Ollama for secure local LLM inference, and MongoDB for state persistence enabled us to build a production-ready architecture blueprint.

This project strengthened our understanding of:

- Multi-agent coordination
- RAG pipelines
- Workflow orchestration
- LLM integration
- State management
- Error handling strategies
- Secure AI system design

## Future Scope & Enhancements

### 1. Add More Specialized Agents

- Compliance Validator Agent
- Fraud Detection Agent
- Quality Assurance Agent

### 2. Scalability Enhancements

- Horizontal scaling with microservices
- Load balancing
- Message queue (Kafka / RabbitMQ)

### 3. Performance Optimization

- Advanced semantic caching
- Async processing
- Database sharding

### 4. Monitoring & Observability

- Distributed tracing
- Centralized log aggregation

### 5. Security Enhancements

- OAuth2 / JWT Authentication
- Role-based Access Control
- End-to-End encryption

### 6. Feature Extensions

- PDF export generation
- Real-time dashboard
- Multi-language policy analysis
- Regulatory auto-updates
- Agent learning with feedback loop